

# From Liquidation Geometry to Price Movement

A Spatial-Density Heatmap and Cascade-Probability Model  
for Cryptocurrency Perpetual Futures

Yigit Onat  
*Independent Researcher*  
yionat@gmail.com

June 2026

## Abstract

Leveraged perpetual-futures positions carry deterministic liquidation prices; when price reaches a dense cluster of them, forced market liquidations impose price impact that can reach the next cluster, producing a self-reinforcing cascade. We formalize this intuition. We construct a *spatial liquidation-density field*—a “liquidity heatmap”—that discretizes the log-return axis around the mid-price into  $K = 200$  bins (10 basis points each,  $\pm 10\%$ ) and maps the USD notional of each tracked position to the bin containing its liquidation price. From this field we extract eight scale- and level-invariant geometric features (nearest-cluster distance, mass concentration, cascade-chain length, directional asymmetry, absorption capacity, cluster velocity, mass-weighted distance, and spatial entropy). We then give a structural price-movement model in which a closed-form domino condition,  $\delta_B - \delta_A < \lambda M_A/m_t$ , determines whether liquidating one cluster triggers the next, with order-book depth acting as a damper (a liquidity spiral) and the directional asymmetry of the field predicting cascade sign. The geometry feeds an online, delayed-target logistic model of cascade probability. The contribution of this paper is the construction and the model, together with an explicit, falsifiable five-gate validation protocol and a candid treatment of the binding practical obstacle: on a decentralized venue the field is observed only through a *sample* of tracked wallets, not the full population of open interest. We formalize this sample-to-population bias, propose an open-interest rescaling that preserves field *shape*, and report the failure mode it produces in practice. We then report first measurements of that bias from a purpose-built Hyperliquid instrument (Section 8), which change what the obstacle *is*. The sampling rate is no longer unknown: a 3,596-wallet registry seeded from a 41,392-wallet leaderboard covers 25–38% of venue open interest per symbol, measured and gate-thresholded on every map. But the measurement also exposes a *second* bias we did not anticipate and which no rescaling can repair: the wallets that hold notional and the wallets that get liquidated are different populations. Of 48 observed liquidations, the registry contained 2.9% of the liquidated *wallets* but 21.9% of the liquidated *notional*—so the field is far better at seeing the mass that cascades than the events that start one. That reading is itself confounded by snapshot timing, and we state the experiment that settles it. We deliberately present no fabricated performance numbers; empirical validation under the stated protocol is the explicit next step.

**Keywords:** market microstructure, liquidation cascade, price impact, perpetual futures, spatial density estimation, fire sales, order-book liquidity, cryptocurrency

**JEL Classification:** G12, G14, G17, C53

## Contents

|          |                                      |          |
|----------|--------------------------------------|----------|
| <b>1</b> | <b>Introduction</b>                  | <b>2</b> |
| 1.1      | The problem . . . . .                | 2        |
| 1.2      | Approach and contributions . . . . . | 3        |

|           |                                                                                  |           |
|-----------|----------------------------------------------------------------------------------|-----------|
| <b>2</b>  | <b>Related Work</b>                                                              | <b>3</b>  |
| <b>3</b>  | <b>The Liquidity Heatmap: Construction</b>                                       | <b>4</b>  |
| 3.1       | Inputs and coordinates . . . . .                                                 | 4         |
| 3.2       | The density field . . . . .                                                      | 4         |
| 3.3       | Derived channels . . . . .                                                       | 4         |
| <b>4</b>  | <b>Spatial Features</b>                                                          | <b>5</b>  |
| <b>5</b>  | <b>The Cascade Mechanism: A Price-Movement Model</b>                             | <b>5</b>  |
| 5.1       | Single-cluster impact . . . . .                                                  | 5         |
| 5.2       | The domino condition . . . . .                                                   | 5         |
| 5.3       | Absorption damping . . . . .                                                     | 6         |
| 5.4       | Direction . . . . .                                                              | 6         |
| <b>6</b>  | <b>A Probabilistic Cascade Model</b>                                             | <b>6</b>  |
| 6.1       | Event label . . . . .                                                            | 6         |
| 6.2       | Features and model . . . . .                                                     | 6         |
| <b>7</b>  | <b>Validation Protocol</b>                                                       | <b>7</b>  |
| <b>8</b>  | <b>Data Availability and the Sample-to-Population Problem</b>                    | <b>7</b>  |
| 8.1       | The sampling rate, measured . . . . .                                            | 8         |
| 8.2       | A second bias: the liquidated population is not the tracked population . . . . . | 9         |
| <b>9</b>  | <b>Discussion</b>                                                                | <b>10</b> |
| <b>10</b> | <b>Conclusion</b>                                                                | <b>10</b> |
|           | <b>References</b>                                                                | <b>11</b> |

# 1 Introduction

## 1.1 The problem

Perpetual futures are the dominant instrument in cryptocurrency derivatives, and they are traded with high leverage. A leveraged position has a deterministic *liquidation price*: the level at which the maintenance-margin constraint binds and the exchange force-closes the position with a market order. These forced orders are price-takers and price-movers. When the mid-price approaches a region of the price axis that is dense with liquidation prices, a move into that region triggers a burst of same-signed market orders whose impact can carry the price into the *next* dense region—a liquidation cascade, the cryptocurrency analogue of a fire sale.

Two facts make this mechanism tractable as a microstructure object. First, the trigger levels are not latent: a position’s liquidation price is a deterministic function of its entry, leverage, and side. Second, the impact channel is the classical one: forced liquidations consume order-book depth, and their price impact is governed by market illiquidity. The open question is therefore not *whether* liquidation geometry matters, but how to *represent* it as a field over the price axis and how that field’s geometry maps to future price movement.

## 1.2 Approach and contributions

We represent the liquidation landscape as a one-dimensional density field  $H(t, \cdot)$  over the log-return axis relative to the current mid-price, updated at 100 ms resolution (Section 3). From the field we extract eight geometric features that are invariant to price level and, being ratios or distances, approximately invariant to the unobserved scale of total open interest (Section 4). We then give a structural model of how the field moves price (Section 5): a single cluster’s impact follows Kyle (1985), a closed-form condition determines whether it triggers the next cluster, order-book depth damps the chain in the manner of Brunnermeier and Pedersen (2009), and the field’s directional asymmetry predicts the sign. The geometry drives an online logistic model of cascade probability with delayed targets (Section 6), gated by an explicit five-criterion validation protocol (Section 7).

Specifically, this paper contributes:

1. A construction of the liquidation-density field in *log-return* coordinates, making bins symmetric in percentage terms and invariant to price level, with an  $O(N_t)$  update.
2. Eight scale-invariant geometric features of the field, each with a microstructural interpretation.
3. A structural price-movement model with a closed-form domino-trigger condition and an order-book absorption damper, linking field geometry to cascade magnitude and direction.
4. An online, delayed-target cascade-probability model and a falsifiable five-gate validation protocol.
5. A formal treatment of the sample-to-population observation problem on decentralized venues, an open-interest rescaling that preserves field shape, and an honest account of the coverage-driven failure mode observed in practice.

We emphasize at the outset that this is a *methods and model* paper. The geometry and the model are fully specified and implemented; a full out-of-sample evaluation under the protocol of Section 7 is stated as future work, and we report no measured performance figures we cannot substantiate.

## 2 Related Work

**Price impact of forced trades.** Kyle (1985) gives the canonical linear price-impact model in which order flow  $q$  moves price by  $\lambda q$ ;  $\lambda$  is the inverse market depth. Forced liquidations are exactly such order flow. Our single-cluster impact and the domino condition are built directly on this primitive, with  $\lambda$  estimated from recent trades in the spirit of Amihud (2002) and Hasbrouck (2009).

**Fire sales and endogenous risk.** The cascade mechanism is a fire sale in the sense of Caccioli, Shrestha, Moore, and Farmer (2014), who analyze contagion in overlapping portfolios, and of Cont and Wagalath (2016), who measure endogenous risk created by feedback between forced selling and price. Our cascade-chain feature is the length of the triggered component of exactly such a contagion graph.

**Liquidity spirals.** Brunnermeier and Pedersen (2009) show how market and funding liquidity interact to amplify shocks in thin markets. Our absorption term—order-book depth divided by cluster mass—is the discrete, level-local form of their damper: a deep book absorbs a liquidation without propagation, a thin one lets price gap through.

**Cryptocurrency liquidations.** Perpetual-futures liquidation cascades are a recognized empirical phenomenon in crypto markets, but they are typically studied ex post from exchange liquidation feeds. Our contribution is a *forward-looking, spatial* representation built from position data, and an explicit treatment of the fact that on a decentralized venue this data is sampled, not complete.

### 3 The Liquidity Heatmap: Construction

#### 3.1 Inputs and coordinates

Let  $m_t$  denote the mid-price at time  $t$ . The system tracks a set of leveraged positions; position  $i$  has notional  $V_i$  (USD), liquidation price  $\ell_i$ , side, and entry price. We place each position on the *log-return axis* relative to the mid,

$$\delta_i = \ln \frac{\ell_i}{m_t}, \quad (1)$$

so that  $\delta_i$  is the (signed) percentage distance of the liquidation level from the current price. The axis is discretized into  $K$  uniform bins on  $[-R, R]$ :

$$\delta_k = -R + (k + \tfrac{1}{2})\Delta, \quad \Delta = \frac{2R}{K}, \quad \mathcal{B}_k = (-R + k\Delta, -R + (k+1)\Delta), \quad k = 0, \dots, K-1, \quad (2)$$

with defaults  $K = 200$ ,  $R = 0.10$  (so  $\Delta = 0.001$ , i.e. 10 basis points per bin, spanning  $\pm 10\%$ ). Working in log-returns makes the bins symmetric in percentage terms and independent of the absolute price level.

#### 3.2 The density field

**Definition 3.1** (Liquidation-density field). The liquidity heatmap is the field

$$H(t, k) = \sum_{i=1}^{N_t} V_i \mathbf{1}[\delta_i \in \mathcal{B}_k] \mathbf{1}[V_i \geq \tau_V], \quad k = \left\lfloor \frac{\delta_i + R}{\Delta} \right\rfloor, \quad (3)$$

i.e.  $H(t, k)$  is the total USD notional whose liquidation price lies in bin  $k$ . Positions with notional below  $\tau_V$  (default \$1,000) or with  $\delta_i \notin [-R, R]$  are discarded.

The field is rebuilt in full each snapshot; each position maps to one bin via a logarithm, an addition, and a floor, so the update is  $O(N_t)$ . There is no temporal smoothing of  $H$  itself; temporal information enters only through the velocity feature below.

#### 3.3 Derived channels

Three further channels give the field local structure (here  $k_0 = K/2$  is the at-the-money bin):

$$\text{cumulative mass: } C(t, k) = \sum_{j \text{ between } k \text{ and } k_0} H(t, j), \quad (4)$$

$$\text{cluster gradient: } G(t, k) = H(t, k) - H(t, k-1), \quad (5)$$

$$\text{temporal change: } T(t, k) = \alpha_1 \frac{H(t, k) - H(t - \Delta t_1, k)}{\Delta t_1} + \alpha_2 \frac{H(t, k) - H(t - \Delta t_2, k)}{\Delta t_2}, \quad (6)$$

with  $\Delta t_1 = 600$  ticks (1 min),  $\Delta t_2 = 3000$  ticks (5 min), and weights  $\alpha_1 = 0.7$ ,  $\alpha_2 = 0.3$ . The cumulative channel measures the fuel between the mid and a level (cascade reach), the gradient locates cluster boundaries, and the temporal channel is the cluster-migration velocity field.

## 4 Spatial Features

Eight scalar features summarize the field’s geometry each snapshot. Write the total mass  $\Sigma_H = \sum_k H(t, k)$ ,  $\bar{H} = \Sigma_H/K$ , and let  $p_k = H(t, k)/\Sigma_H$  be the normalized mass. Let  $\tau_{cl} = \$10^6$  and  $\tau_{ch} = \$10^5$  be the cluster and chain thresholds.

| Feature                     | Definition                                                                                                                            | Interpretation                                                       |
|-----------------------------|---------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------|
| hm_nearest_cluster_distance | $d_{\min} = \min\{ \delta_k  : H(t, k) > \tau_{cl}\}$                                                                                 | distance to the nearest large cluster (smaller = cascade closer)     |
| hm_cluster_mass_ratio       | $CMR = H(t, k^*)/\bar{H}$ , $k^* = \arg \max_k H$                                                                                     | peak concentration relative to mean                                  |
| hm_cascade_chain_length     | longest run of consecutive bins with $H > \tau_{ch}$ within $ \delta  \leq 0.05$                                                      | domino potential (triggered-chain length)                            |
| hm_asymmetric_cascade       | $ACP = \frac{\sum_{0 < \delta \leq .05} H - \sum_{-.05 \leq \delta < 0} H}{\sum_{ \delta  \leq .05} H + \epsilon}$                    | directional bias: $> 0$ short-squeeze, $< 0$ long-cascade            |
| hm_absorption_capacity      | $A = \min\left(\frac{D^{\text{bid}}(.02)}{\text{down mass} + \epsilon}, \frac{D^{\text{ask}}(.02)}{\text{up mass} + \epsilon}\right)$ | book depth per unit liquidation mass ( $< 1$ = insufficient damping) |
| hm_cluster_velocity         | $v = (d_{\min}(t) - d_{\min}(t - \tau))/\tau$ , $\tau = 600$                                                                          | rate of approach of the nearest cluster ( $< 0$ approaching)         |
| hm_mass_weighted_distance   | $\bar{d} = \frac{\sum_k  \delta_k  H(t, k)}{\Sigma_H + \epsilon}$                                                                     | centre of liquidation gravity                                        |
| hm_heatmap_entropy          | $S = -\sum_k p_k \ln p_k$                                                                                                             | spatial concentration ( $S$ low = tight clusters)                    |

Table 1: The eight liquidity-heatmap features.  $\epsilon$  is a small constant (1 USD). Absorption uses order-book depth within 2% on each side; when a side carries no mass its ratio is treated as  $+\infty$ , and  $A$  is clamped to  $[0, 100]$ .

These features are deliberately constructed to be *ratios or signed distances* (Table 1): CMR, ACP,  $S$ ,  $\bar{d}$ , and  $d_{\min}$  are unchanged if the entire field is rescaled by a common factor. This is the property that makes them usable when only a fraction of total open interest is observed (Section 8).

## 5 The Cascade Mechanism: A Price-Movement Model

### 5.1 Single-cluster impact

Let cluster  $A$  at  $\delta_A$  carry mass  $M_A = H(t, k_A)$ . Treating its forced liquidation as an order of notional  $M_A$  and applying Kyle (1985) with depth parameter  $\lambda$  (estimated from recent trades), its price impact in log-return units is

$$\Delta p_A = \lambda \frac{M_A}{m_t}. \quad (7)$$

### 5.2 The domino condition

A second cluster  $B$  at  $\delta_B > \delta_A$  (on the same side of the mid) is reached if and only if cluster  $A$ ’s impact closes the gap:

$$\boxed{\delta_B - \delta_A < \lambda \frac{M_A}{m_t}} \quad (8)$$

Equation (8) is the structural heart of the model. It depends on the *spacing* of clusters (captured by the gradient channel and the chain-length feature) and their *mass* (the concentration feature), mediated by illiquidity  $\lambda$ . When it holds across a run of clusters  $\delta_1 < \dots < \delta_n$ , the total move is the sum of triggered masses,

$$\Delta p_{\text{total}} \approx \frac{\lambda}{m_t} \sum_{j: \text{triggered}} M_j, \quad (9)$$

and the number of triggered nodes is exactly the cascade-chain feature, read off a Caccioli et al. (2014) contagion graph.

### 5.3 Absorption damping

Order-book depth opposes the cascade. With bid depth  $D^{\text{bid}}(\delta_A)$  available to absorb a downward liquidation at  $A$ , define the damping

$$\eta = \min\left(1, \frac{D^{\text{bid}}(\delta_A)}{M_A}\right), \quad \Delta p_A^{\text{eff}} = (1 - \eta) \Delta p_A. \quad (10)$$

When  $\eta \rightarrow 1$  (a deep “wall,” i.e. absorption  $A \gg 1$ ) the impact is absorbed and the chain breaks; when  $\eta \rightarrow 0$  (a thin book) it is amplified—the Brunnermeier–Pedersen (2009) spiral. This is the mechanism behind support/resistance walls and absorption.

**Proposition 5.1** (Magnet and rejection). Under (8)–(10), a large cluster within reach acts as a price *magnet*: conditional on the mid entering its neighborhood, the probability of a move toward it increases with  $M_A$  and decreases with the local absorption  $A$ . Conversely, a cluster fronted by a deep book ( $\eta \approx 1$ ) produces *rejection* rather than breakout.

### 5.4 Direction

The sign of the field’s near-money asymmetry predicts the cascade direction. With ACP as in Table 1, mass concentrated above the mid is short-liquidation fuel (an upward squeeze) and mass below is long-liquidation fuel (a downward cascade), so the predicted direction is  $\text{sign}(\text{ACP})$ .

## 6 A Probabilistic Cascade Model

### 6.1 Event label

We define a cascade over horizon  $T$  as a large *liquidation-driven* move—a price move conjoined with abnormal liquidation volume, which distinguishes it from an ordinary news move:

$$y_t = \mathbf{1}\left[|\Delta p_{t \rightarrow t+T}| > X\% \wedge \text{LiqVol}_{t \rightarrow t+T} > Y\right], \quad \Delta p = \ln \frac{m_{t+T}}{m_t}, \quad (11)$$

with defaults  $X = 3\%$ ,  $Y = \$500\text{K}$ ,  $T = 3000$  ticks (5 min), giving a base rate near 5%. Because per-tick liquidation volume is not directly observable on the venue, the implementation currently uses the price-threshold component as a proxy—a conservative choice we flag explicitly.

### 6.2 Features and model

The model consumes an 11-dimensional vector: the eight features of Table 1 plus three interactions chosen from the mechanism,

$$\mathbf{x}(t) = \left[ d_{\min}, \text{CMR}, L, \text{ACP}, A^{-1}, v, \bar{d}, S, \underbrace{L \cdot A^{-1}}_{\text{fuel, no damper}}, \underbrace{v \cdot \mathcal{I}}_{\text{approach} \times \text{illiquidity}}, \underbrace{\text{ACP} \cdot \text{sign}(f)}_{\text{aligned with funding}} \right]^\top, \quad (12)$$

where  $L$  is the chain length,  $\mathcal{I}$  a composite illiquidity, and  $f$  the funding rate. Cascade probability is an online logistic regression,

$$P(\text{cascade} \mid \mathbf{x}) = \sigma(\beta_0 + \beta^\top \mathbf{x}), \quad \sigma(z) = \frac{1}{1 + e^{-z}}, \quad (13)$$

trained by stochastic gradient descent on the cross-entropy loss with  $\ell_2$  regularization ( $\lambda_{\text{reg}} = 10^{-3}$ ). With per-step gradient  $g_t = \hat{p}_t - y_t$ ,

$$\beta \leftarrow \beta - \eta_t(g_t \mathbf{x}_t + \lambda_{\text{reg}} \beta), \quad \beta_0 \leftarrow \beta_0 - \eta_t g_t, \quad \eta_t = \eta_0 \gamma^t, \quad (14)$$

with  $\eta_0 = 0.01$ ,  $\gamma = 0.9999$ . Because the label (11) is only known  $T$  ticks later, updates are applied from a *delayed-target* buffer once each observation matures; inputs are standardized online by Welford’s algorithm, and the logit is clipped to  $[-20, 20]$ . A logistic model is deliberate: cascades are rare ( $\approx 5\%$ ), leverage regimes are non-stationary, and the coefficients  $\beta$  are interpretable as the weight the data place on each geometric feature.

## 7 Validation Protocol

The model is only credible if it clears pre-registered gates. We specify five; they are *thresholds to be met*, not results claimed here.

| Gate                 | Criterion                                                                                       |
|----------------------|-------------------------------------------------------------------------------------------------|
| G1 — Discrimination  | out-of-sample AUC $> 0.65$ and top-decile lift $> 2\times$                                      |
| G2 — Net information | $\text{IC}_{\text{net}} = \text{IC}_{\text{gross}} - \frac{2\bar{s}}{\sigma_r \sqrt{T}} > 0.02$ |
| G3 — Stability       | walk-forward median AUC $> 0.60$ , $< 20\%$ of windows below 0.55                               |
| G4 — Universality    | BTC, ETH, SOL each pass G1, and $\cos(\beta_i, \beta_j) > 0.6$ pairwise                         |
| G5 — Orthogonality   | $R^2 < 0.30$ regressing $\hat{p}$ on a spread/depth composite                                   |

Here  $\bar{s}$  is the cascade-event spread (a cost proxy) and  $\sigma_r$  the return volatility. G2 ensures the signal survives transaction costs; G5 ensures it is not a relabeling of ordinary liquidity. The prior empirical anchor that motivates the whole construction is the hypothesis-test result that coarse liquidation-cluster proximity predicts large moves with lift  $> 2\times$ ; the heatmap’s claim is that its eight *spatial* features carry information *incremental* to that coarse proximity—a claim Gate G5 and an explicit ablation against `hm_nearest_cluster_dist` alone are designed to test.

## 8 Data Availability and the Sample-to-Population Problem

The model’s binding obstacle is observational, not theoretical. On a decentralized venue, position data is obtained by querying the clearinghouse state of *known wallet addresses*; no endpoint returns the full population of open positions. Wallets are discovered from the public trade stream and polled, yielding coverage on the order of 50–500 wallets per symbol—an *unknown* fraction of open interest. The field of Definition 3.1 is therefore a sample,  $H^{\text{obs}}(t, k) = \rho_t H(t, k) + \text{noise}$ , with sampling rate  $\rho_t$  unobserved and time-varying.

That was the state of the world when this section was first written. It has since been measured, by a purpose-built Hyperliquid instrument whose design and gate protocol are the subject of a companion paper. Subsections 8.1 and 8.2 report what the measurement found; the theory in between (what survives sampling, and the rescaling) is unchanged by it.

**What survives sampling.** The scale-invariant features— $d_{\text{min}}$ , CMR, ACP,  $\bar{d}$ ,  $S$ —depend on the field’s *shape*, which is preserved in expectation under uniform sampling; and the largest, cascade-causing positions are disproportionately likely to be discovered (they dominate trade flow). What is biased are the *absolute-mass* quantities: absorption capacity  $A$  is overestimated (observed mass underestimates the true denominator), and chain length may be undercounted when untracked positions fill gaps in (8).

**A shape-preserving rescaling.** Because total open interest *is* observable, we can restore scale without distorting shape,

$$H^{\text{scaled}}(t, k) = H^{\text{obs}}(t, k) \cdot \frac{\text{OI}_{\text{total}}(t)}{\sum_k H^{\text{obs}}(t, k)}, \quad (15)$$

which leaves every ratio feature unchanged while correcting the mass-dependent ones. Validating (15) against realized cascade magnitudes is itself a clean empirical sub-study.

**Observed failure mode.** We report candidly that, in a preliminary multi-symbol backtest, the cascade-probability algorithm produced *zero* trades and the nearest-cluster-distance feature frequently returned its no-data sentinel—direct evidence that, under the wallet coverage available at the time, the observed field was too sparse for the absolute-mass features to fire. This was the single most important practical result of the work to that point: the geometry and the model are sound, but their utility is gated on coverage, and the first priority is therefore data—aggressive wallet discovery and the rescaling (15)—not further modeling.

## 8.1 The sampling rate, measured

Acting on that priority produces the first direct estimate of  $\rho_t$ . Hyperliquid publishes a leaderboard exposing 41,392 wallets with equity and windowed volume—roughly forty times the population the trade-stream discovery of the previous paragraph reaches. Seeding a registry from it is therefore a different regime, and the two natural seeds behave differently (probe of 2026-08-07): the top 2,000 wallets *by equity* span 61% of venue open interest but only 28% of them hold any position at a given moment, while the top 2,000 *by weekly volume* span 39% with 46% active. The union reaches 69% of venue OI (BTC 68%, ETH 79%, HYPE 70%, SOL 66%). Mass and activity are thus carried by distinguishable populations—a fact we return to in Section 8.2, where it becomes the central obstacle rather than a seeding detail.

The operational registry holds 3,596 wallets, and its cost sets the map’s time resolution: one full sweep of clearinghouse state takes approximately five minutes and consumes the entire per-IP request-weight budget, leaving nothing for the trade tape. Polling frequency would therefore *become* the field’s resolution, which is why positions are maintained between snapshots by reconstruction from the public fill stream and polling is demoted to reconciliation.

Coverage is then computed per map and per symbol as

$$\hat{\rho}_t = \frac{\sum_k H^{\text{obs}}(t, k)}{s \cdot \text{OI}_{\text{total}}(t)}, \quad s = 2, \quad (16)$$

where  $s$  is the side convention: an exchange that counts each contract once in  $\text{OI}_{\text{total}}$  reports half the position notional actually outstanding, since every long has a matching short and a registry sees both. Taking  $s = 2$  is the conservative reading and halves the headline number; it is printed on every map alongside the coverage figure, because being wrong about  $s$  is a factor of two on the one quantity by which the field is judged. Under (16), gate verdicts recorded on 2026-08-08 give

$$\hat{\rho}_{\text{BTC}} = 0.315, \quad \hat{\rho}_{\text{ETH}} = 0.380, \quad \hat{\rho}_{\text{SOL}} = 0.247,$$

against a 0.25 minimum below which the map’s gate refuses to pass. (The gap between these figures and the 69% scoping probe above is very nearly the factor  $s$  itself; whether the probe was computed one-sided is an open verification item, and we report both numbers rather than silently adopting the flattering one.)

Two consequences for this paper. First,  $\rho_t$  is no longer an unobserved nuisance parameter: it is a measured, logged series, so the rescaling (15) can be applied with a known multiplier and audited after the fact rather than assumed. Second, a related fidelity question becomes measurable. Where the venue publishes its own liquidation price per position, that published

value is the source of truth and the field needs no model; a derived price is needed only for positions whose size changed since the last snapshot. Diffing the derivation against the published field over 2,926 positions gives a median relative error of 0.152 and agreement to within  $10^{-4}$  for 40.8% of positions (58% among accounts holding no isolated margin, over 1,074 cross positions). The residual is unexplained—size-tiered maintenance margin is the leading suspect—and is the reason the derivation is used only for maintenance and never as the primary source.

## 8.2 A second bias: the liquidated population is not the tracked population

Measuring  $\rho_t$  was expected to be the end of the observational problem. It was not. The measurement exposes a second bias, distinct from the first in kind, and one that no mass rescaling can repair.

Equation (15) corrects a *magnitude* error under the assumption that the observed field is a uniform thinning of the true one. That assumption is what fails. The registry is seeded by equity because the map wants whoever holds size; but the wallets that actually get liquidated are small and highly leveraged. Sampling proportional to notional is precisely *not* sampling proportional to liquidation events.

The magnitudes are stark. Of 48 liquidations observed on Hyperliquid and appended to the instrument’s ledger, the registry contained

2.9% of the liquidated *wallets*, but 21.9% of the liquidated *notional*

(\$223,491 of realized liquidation notional across 35 distinct wallets; a second reading a cycle later gave 2.8% and 21.9% over 49 events). The order-of-magnitude gap between the two percentages *is* the second bias, stated quantitatively. It is also, for this paper, a qualified reprieve: the field is built to predict cascades, cascades are driven by notional, and a fifth of liquidated notional is not nothing. What the field is poor at is seeing the long tail of small forced exits—which matters if those exits are what *initiates* a chain in (8).

**Why the reading cannot yet be believed.** Both readings show 0.0% of liquidated notional *mapped*—i.e. attributable to a position the field had marked in advance—and that zero is confounded rather than informative. A liquidated wallet no longer holds the position, so a registry snapshot taken *after* a liquidation shows nothing that could have predicted it. Until snapshots reliably precede the events they are scored against, the measurement cannot distinguish “different populations” from “we looked too late”. Reporting the confound rather than the headline is the whole point: the 0.0% would otherwise read as a refutation of the model, when it is currently a statement about observation order.

**The experiment that settles it.** A scheduled snapshot-then-observe cycle—sweep the registry, then scan for liquidations, appending an overlap reading each pass—turns the single confounded reading into a series, and the series discriminates:

- mapped notional fraction climbing off zero  $\Rightarrow$  per-position scoring works and the field can be validated as constructed;
- mapped fraction pinned at zero while the notional fraction holds near 22%  $\Rightarrow$  the registry sees the right wallets but never *before* they are liquidated, and validation must move from per-position to cluster-level scoring—did liquidations land in the price buckets the field marked, regardless of whose they were;
- the notional fraction collapsing as well  $\Rightarrow$  neither, and the registry needs a third seed drawn by *liquidation risk* (margin ratio, distance to liquidation price) rather than by size or activity.

Only the first outcome validates the instrument as specified; the second demotes the claim this paper can make from position-level to cluster-level, which is weaker but still falsifiable; the third reopens registry construction. None of the three is available from analysis—only from calendar time under capture. Correspondingly, the instrument’s map gate currently records **FAIL** on its predictive criterion, on the grounds that a criterion which cannot yet be evaluated must not be allowed to pass by default.

**What this does to the validation protocol.** Section 7’s protocol is unchanged in structure but acquires a precondition: G1–G5 are not merely unmet, they are *unevaluable* until the cycle above resolves. We regard stating that explicitly as more useful than reporting the protocol as pending.

## 9 Discussion

The model makes several falsifiable claims that a venue with adequate coverage could test directly. The domino condition (8) predicts that the probability of a multi-cluster cascade jumps discontinuously as inter-cluster spacing falls below  $\lambda M_A/m_t$ ; the absorption relation (10) predicts that deep near-money book depth converts approach into rejection rather than breakout (Proposition 5.1); and the directional claim predicts that  $\text{sign}(\text{ACP})$  leads cascade sign. Each is a clean hypothesis with a pre-specified statistic. The most important open question is incremental information: do the eight spatial features add predictive content beyond simple cluster proximity? The construction is designed so that this question is answerable by ablation, and Gate G5 guards against the field being a disguised liquidity proxy.

The honest status is that this is a fully specified instrument awaiting the data to validate it. That ordering—specify the mechanism, expose the observational constraint, fix the data, then measure—is, we argue, the correct one for a phenomenon whose theoretical structure is clear but whose measurement on a decentralized venue is genuinely hard.

Section 8 is the first return from executing that ordering, and it is worth being precise about what it did and did not settle. It settled the constraint the paper originally named: the sampling rate is measured, logged, and gate-thresholded rather than unknown. It did not clear the path, because the act of measuring surfaced a bias of a different kind—mass and events are carried by different wallets—which the rescaling (15) is structurally unable to correct, since that correction assumes uniform thinning. The useful lesson is that “insufficient coverage” was an incomplete diagnosis: coverage *of what* turns out to be the operative question, and a field can be simultaneously adequate for the mass that sustains a cascade and inadequate for the events that start one.

## 10 Conclusion

We have constructed a spatial liquidation-density field for cryptocurrency perpetual futures, extracted eight scale-invariant geometric features from it, and given a structural price-movement model in which a closed-form domino condition, an order-book damper, and a directional asymmetry map the field’s geometry to cascade magnitude, occurrence, and sign. The geometry drives an online, delayed-target cascade-probability model, gated by an explicit five-criterion protocol. We have been deliberate in not reporting performance numbers we cannot substantiate, and in foregrounding the binding constraint: on a decentralized venue the field is sampled, not observed in full, and the immediate priority is the coverage and rescaling needed to make the absolute-mass features fire. With that data in hand, every claim above becomes a pre-registered, testable hypothesis.

Section 8 reports the first measurements against that constraint. The sampling rate is now known—25–38% of venue open interest per symbol, from a 3,596-wallet registry, logged per

map and gate-thresholded—and the derived liquidation price is calibrated against the venue’s published one (40.8% exact over 2,926 positions). Both were open unknowns in the first version of this paper. In their place stands a sharper and more interesting obstacle: the tracked population is selected on notional and the liquidated population is selected on leverage, and the two barely intersect at the wallet level (2.9%) while overlapping substantially at the notional level (21.9%). Whether that asymmetry is fatal to position-level validation or merely demotes it to cluster-level is not decidable from the data in hand; it is decidable from a scheduled snapshot-then-observe series, which is running. We report the confound rather than the headline because the headline—zero mapped notional—is currently an artifact of observation order and would be read as a refutation it has not earned.

## References

1. Amihud, Y. (2002). Illiquidity and stock returns: cross-section and time-series effects. *Journal of Financial Markets*, 5(1), 31–56.
2. Brunnermeier, M.K. & Pedersen, L.H. (2009). Market liquidity and funding liquidity. *Review of Financial Studies*, 22(6), 2201–2238.
3. Caccioli, F., Shrestha, M., Moore, C. & Farmer, J.D. (2014). Stability analysis of financial contagion due to overlapping portfolios. *Journal of Banking & Finance*, 46, 233–245.
4. Cont, R. & Wagalath, L. (2016). Fire sales forensics: measuring endogenous risk. *Mathematical Finance*, 26(4), 835–866.
5. Hasbrouck, J. (2009). Trading costs and returns for U.S. equities: estimating effective costs from daily data. *Journal of Finance*, 64(3), 1445–1477.
6. Kyle, A.S. (1985). Continuous auctions and insider trading. *Econometrica*, 53(6), 1315–1335.
7. Roll, R. (1984). A simple implicit measure of the effective bid-ask spread in an efficient market. *Journal of Finance*, 39(4), 1127–1139.
8. Bandt, C. & Pompe, B. (2002). Permutation entropy: a natural complexity measure for time series. *Physical Review Letters*, 88(17), 174102.
9. Welford, B.P. (1962). Note on a method for calculating corrected sums of squares and products. *Technometrics*, 4(3), 419–420.